

Playbook-IQ — MVP Roadmap

Misk Launchpad Cohort 10 · Assignment #2 · 28 September – 30 October 2026

Product stage: Prototype **Sprint cadence:** 1 week · **Number of sprints:** 5 **Team:** Founder (full-time)
+ Data Operations Specialist (part-time)

The pilot — four clubs, three matches

One match analysed produces two team reports. That is the whole design: three matches, five reports, **four different clubs**, for the work of three.

MATCH	DATE	REPORTS	PURPOSE
Rehearsal	31 Aug (already played)	Al-Ittifaq + opponent	No deadline. Learn what a report costs and whether it is any good
Live 1	12 Oct	Al-Ittifaq + opponent	Delivered within 48h
Live 2	18 Oct	Al-Ittifaq + opponent	Delivered within 48h

Al-Ittifaq U21 is the anchor — top of the table, three points clear, access through their former U18/U21 analyst who is now with the first team. They receive three reports, which also demonstrates the development tier without us having to sell it.

The three opponent clubs are cold doors, opened with a finished report. We produce the analysis first and approach with it in hand. A real report gets past a gatekeeper that a message cannot. Routes in: Misk, and Salman at the Ministry of Sport. **If a club will not take the meeting, that is a finding, not a failure** — it tells us what reaching this buyer actually costs.

Why one match per club and not four. Per mentor guidance: one match is enough to test whether the analysis is worth anything. Four matches tests whether *development tracking over time* works — which is a **later service tier**, not the thing to validate first. Four independent clubs reacting beats one friendly club reacting four times.

What wins is the data and the metrics, not the volume. Fewer matches, better analysis in each.

Follow-through: Al-Ittifaq's next fixture is **2 November**, after this window. Phase 2 story.

How success is measured

After each report, the same three questions to whoever reads it:

1. Which findings did you already know?
2. Which were new to you?
3. Which changed something you did?

METRIC	TARGET BY 30 OCT
Clubs reached — report delivered and read	4 of 4
Actions taken — findings that changed a decision	≥ 1 per club that engages
Novelty rate — findings rated new ÷ findings delivered	≥ 40%
Turnaround — match end to report delivered	< 48 h on both live matches
Human hours per match	measured every match, falling

Kill criteria

Fewer than 2 of the 4 clubs ask for a second report, and none accepts any price, by 30 October → we stop targeting club development squads and reconsider the segment.

Three tracks

TRACK	BUYS US
A — Pilot	Evidence. The willingness-to-pay answer we do not have
B — Models	Cost. Every point of detection recall reduces manual identity work
C — Benchmark	Proof. A number comparable to published research

Sprint plan

#	DATES	A — PILOT	B — MODELS	C — BENCHMARK
1	28 Sep – 4 Oct	Report template agreed with the Al-Ittifaq analyst · target clubs and routes in identified	Compute secured · CVAT annotations migrated · detector retraining started	SoccerNet-GSR obtained, reference baseline reproduced
2	5 – 11 Oct	Rehearsal: 31 Aug match, both team reports. Hours logged, analyst critique captured	Retrained detector live · recall and identity-break rate measured against the 952-fragment baseline	Our footage converted to ground truth with verified identities
3	12 – 18 Oct	Match 12 Oct: both reports within 48h. Al-Ittifaq delivered; first cold door knocked	Pitch model validated on pilot footage	Pipeline scored on the GSR benchmark — first comparable number
4	19 – 25 Oct	Match 18 Oct: both reports within 48h. Second cold door knocked. Reactions collected	Review tooling in use · identity work per match measurably reduced	Failure modes ranked
5	26 – 30 Oct	Four clubs' reactions compiled · a price named and the reaction recorded	Human-hours per match reported as a measured number	Results written up and taken to a research group

Every sprint ends in a measured result or something delivered to a real user.

Running through all five: reaching the budget holder at the club through the staff, and a PDPL consent note for under-18 player data.

The report template — design principles

Agreed in Sprint 1, before the first report is produced. These are framing choices, not compliance work, and they are expensive to unwind once a club has seen version one.

No within-squad leaderboard in the development report. Ranking a club's own players 1-to-N invites a cut list; a development profile invites a coaching conversation. Same data, different consequence. Show each player against their own previous form and against positional norms — not against their team-mates in a table.

Comparison is the point in the recruitment report. A separate report type, for management. Evaluating a player you might sign requires comparing him to others, and that is the tier management

actually pays for. The subject there is someone you might buy, not someone you might release — so the leaderboard logic that is wrong in one product is the whole product in the other.

Profiles, not rankings. A player is not objectively better or worse — he is a different profile, suited to a different role, style and budget. Framing comparison as *fit* rather than *quality* is both more honest and more sellable, and it is how recruitment actually works.

Every report states its basis. What it measures, what it does not, and the sample it rests on (one match, in this pilot). Our pipeline already carries a confidence value on every output row — surfacing that is a differentiator, not a weakness.

Supporting buy and sell decisions is deliberate. It is a primary reason the product exists and the highest-value thing management buys. It also raises the stakes of being wrong, so it needs written agreements with clubs covering scope, use and liability before any report is used that way commercially. **Worth asking whether Misk or the Ministry of Sport can help with a standard framework** — building that institutional layer is part of why they back startups in this sector.

Role-based access — a named later tier, not this quarter

The platform will eventually serve three audiences from the same data, with different views:

ROLE	SEES
Management	Cross-market comparison, profiles against budget and role. The recruitment tier
Coaching staff	Squad development, session-level findings. No internal leaderboard
Players	Their own trend, and their standing against anonymised positional benchmarks — enough to understand their market value and compete, without a ranking of their own dressing room

The player view is the most valuable second revenue line (players and agents pay directly) and the highest-risk surface. Named internal comparison is where dressing rooms break, and it is where the minors question bites hardest — so the limit is: **own trend always, anonymised benchmarks yes, named team-mates no**, and named comparison only where that data is already public.

This is on the not-list for the pilot. Three role-based views is a quarter of work and would eat the five weeks. Recording it here so it is a decision, not an omission.

Main dependency — compute

Training runs on Roboflow (plan upgrade, Sprint 1). The pilot needs inference over **4 matches × 90 min × 25 fps ≈ 540,000 frames**, on the order of **10–25 GPU-hours** across the five weeks. This gates

Sprints 1–4.

Routes: Roboflow hosted batch processing, cloud startup credits, the programme, a research partner.

MVP Scope Canvas

01 One user	The U21 analyst or coach at a Saudi professional club — starting with Al-Ittifaq, whose analyst reads our reports
02 One job	<i>"I want to know before the next session which of my players is actually developing, and which isn't."</i>
03 Riskiest assumption	That a club will pay for this rather than keep relying on the analyst's eye. Nobody has paid us anything
04 Smallest killing test	Three matches, five reports, four clubs — produced by hand and delivered unsolicited. No automation, no product
05 Kill number and date	Fewer than 2 of 4 clubs ask for a second report, and none accepts any price, by 30 October → stop
06 The not-list	No tactical camera or hardware · no opposition scouting · no first-team product · no mobile app · no automated re-identification · no multi-match development tracking and no role-based player or management portals — those are the next tiers, not the MVP. Not this quarter

Key planned features

FEATURE	SPRINT
Player profiles in the prototype	shipped
Retrained player detector (RF-DETR)	1–2
Report template — development and recruitment variants	1
Match, team and player analysis on pilot footage	1–4
Arabic assistant answering the coach's own questions	3
Tracklet review tooling (jersey-number assisted)	3–4
Cross-match comparison for the anchor club	4
GSR benchmark score on our own footage	3–4
PDPL consent note for minors	4

Beyond 30 October

WHAT	
Automatic identity	Open research problem with a public benchmark. Pursued with a research partner. The outcome is must-have; automating it is should-have
Full automation	Manual capacity is roughly three squads. Automation becomes urgent at customer #4
The research network	Specialist partners by domain — computer vision, tactical modelling, sports science — with Playbook-IQ holding the club relationships, deployment and product. <i>"Football by the book."</i>
